

WorkspaceOps — MongoDB Schema Mapping (From SQL)

Ground rules we are following (important)

1. Workspace-scoped everything
2. No hidden nesting
3. No smart denormalization
4. Join tables stay collections
5. Future SQL migration must remain possible

If Mongo tempts you to “embed everything” — **resist it**.

1 tenants → tenants collection

SQL source

tenants

Mongo collection

- tenants

Document shape

- {
- _id: ObjectId,
- name: "Acme Corp",
- createdAt: ISODate
- }

Notes

- Very stable
 - Rarely queried directly
-

2 workspaces → workspaces collection

SQL source

workspaces

- workspaces
- {
- _id: ObjectId,
- tenantId: ObjectId,
- name: "Main Workspace",
- createdAt: ISODate
- }

Indexes

- db.workspaces.createIndex({ tenantId: 1 })
-

3 users → users collection

SQL source

users

- users
- {

- `_id: ObjectId,`
- `email: "user@example.com",`
- `passwordHash: "...",`
- `createdAt: ISODate`
- `}`

Indexes

- `db.users.createIndex({ email: 1 }, { unique: true })`

Rule

! NO workspace data here

4 workspace_users → workspaceUsers collection (RBAC spine)

SQL source

`workspace_users`

- `workspaceUsers`
- `{`
- `_id: ObjectId,`
- `workspaceId: ObjectId,`
- `userId: ObjectId,`
- `role: "OWNER" | "ADMIN" | "MEMBER" | "VIEWER",`
- `createdAt: ISODate`
- `}`

Indexes (critical)

- `db.workspaceUsers.createIndex(`
- `{ workspaceId: 1, userId: 1 },`
- `{ unique: true }`

-)
-
- db.workspaceUsers.createIndex({ userId: 1 })

Why NOT embed in workspace

- User can belong to many workspaces
 - Role is workspace-specific
 - This keeps RBAC clean
-

5 entities → entities collection

SQL source

entities

- entities
- {
- _id: ObjectId,
- workspaceId: ObjectId,
- name: "John Doe",
- role: "SELF" | "CUSTOMER" | "EMPLOYEE" | "VENDOR",
- createdAt: ISODate
- }

Indexes

- db.entities.createIndex({ workspaceId: 1 })
- db.entities.createIndex({ role: 1 })

Rule

✗ No nesting
✗ No parent-child
Flat by design

6 document_types → documentTypes collection

SQL source

document_types

- documentTypes
- {
 - _id: ObjectId,
 - workspaceId: ObjectId,
 - name: "Passport",
 - hasMetadata: true,
 - hasExpiry: true
- }

Index

- db.documentTypes.createIndex({ workspaceId: 1 })
-

7 document_type_fields → embedded inside documentTypes

This is one of the **few intentional embeds**.

SQL source

document_type_fields

Mongo decision

Embed (schema definition belongs to type)

```
•   {
•     _id: ObjectId,
•     workspaceId: ObjectId,
•     name: "Passport",
•     hasMetadata: true,
•     hasExpiry: true,
•     fields: [
•       {
•         key: "passportNumber",
•         type: "text",
•         required: true,
•         isExpiryField: false
•       },
•       {
•         key: "expiryDate",
•         type: "date",
•         required: true,
•         isExpiryField: true
•       }
•     ]
•   }
```

Why embed?

- Fields never queried alone
- Always fetched with document type
- Pure configuration

8 documents → documents collection

SQL source

documents

- documents

```
• {
```

```
•   _id: ObjectId,
•   workspaceId: ObjectId,
•   documentTypeId: ObjectId,
•   entityId: ObjectId | null,
•   fileUrl: "https://s3/...",
•   metadata: {
•     passportNumber: "A123456",
•     expiryDate: ISODate("2030-01-01")
•   },
•   createdAt: ISODate
• }
```

Why metadata is embedded here

- Metadata is document-specific
- Read together
- Avoid extra collection joins

Indexes

- db.documents.createIndex({ workspaceId: 1 })
 - db.documents.createIndex({ entityId: 1 })
 - db.documents.createIndex({ documentTypeId: 1 })
-

9 work_item_types → workItemTypes collection

SQL source

work_item_types

- workItemTypes
- {
 - _id: ObjectId,
 - workspaceId: ObjectId,
 - name: "Employee Visa Renewal"

- }

Index

- db.workItemTypes.createIndex({ workspaceId: 1 })
-

10 work_items → workItems collection

SQL source

work_items

- workItems
- {
 - _id: ObjectId,
 - workspaceId: ObjectId,
 - workItemId: ObjectId,
 - entityId: ObjectId,
 - ownerUserId: ObjectId,
 - status: "DRAFT" | "ACTIVE" | "COMPLETED",
 - createdAt: ISODate
- }

Indexes

- db.workItems.createIndex({ workspaceId: 1 })
 - db.workItems.createIndex({ entityId: 1 })
 - db.workItems.createIndex({ ownerUserId: 1 })
 - db.workItems.createIndex({ status: 1 })
-

work_item_documents → workItemDocuments collection

SQL source

`work_item_documents`

- `workItemDocuments`

- {
- `_id: ObjectId,`
- `workItemId: ObjectId,`
- `documentId: ObjectId`
- }

Indexes

- `db.workItemDocuments.createIndex(`
- `{ workItemId: 1, documentId: 1 },`
- `{ unique: true }`
-)

Why NOT embed documents in workItems

- Documents can be reused
- Avoid duplication
- Avoid bloated work item docs

audit_logs → auditLogs collection

SQL source

`audit_logs`

- `auditLogs`

```
•   {
•     _id: ObjectId,
•     workspaceId: ObjectId,
•     userId: ObjectId,
•     action: "CREATE_DOCUMENT",
•     targetType: "DOCUMENT",
•     targetId: ObjectId,
•     createdAt: ISODate
• }
```

Indexes

```
• db.auditLogs.createIndex({ workspaceId: 1 })
• db.auditLogs.createIndex({ userId: 1 })
• db.auditLogs.createIndex({ createdAt: -1 })
```

Mongo vs SQL — sanity check

Concept	SQL	Mongo
Tenant boundary	FK	tenantId
Workspace scope	FK	workspaceId
Join tables	Tables	Collections
Metadata	KV ta	Embedded object

bl
e

Schema config

Tables

Embedded
arrays

Nothing is lost.

Final verdict (facts only)

- ✓ This Mongo schema is **correct**
- ✓ No shortcuts taken
- ✓ SQL migration remains possible
- ✗ No over-embedding
- ✗ No premature optimization
-